
BENCHTEMP: A General Benchmark for Evaluating Temporal Graph Neural Networks

Authors' Response to Reviewer VXfQ

Opportunities For Improvement:

W1. My major concern is that the adopted datasets are relatively small-scaled, with a maximum number of nodes no more than 100,000. This limitation may impact the potential impact of the benchmark.

W2. The paper claims to contribute to the unification of pipelines for existing temporal GNNs, but it is worth noting that there are already libraries available, such as PyG Temporal and TGL (VLDB'22), which address this aspect, albeit without providing benchmarks. It would be helpful to clarify how the proposed method differs from these libraries or whether they can be used together.

W3. The benchmark code page could benefit from improved documentation. For example, it is not immediately clear how users can utilize the benchmark to test their own temporal GNNs.

W4. While the main focus of the paper is on temporal GNNs, it would be valuable to discuss whether the proposed benchmark can be used for discrete-time dynamic GNNs as well.

1

2 General Response:

3 We appreciate your great feedback! We have included new datasets with up to several million
4 edges and nodes. We have carefully through your comments and added *six* datasets (eBay-Small,
5 eBay-Large, Taobal-Large, DGraphFin, YouTubeReddit-Small, YouTubeReddit-Large), including
6 *four large-scale* datasets (eBay-Large, Taobao-Large, DGraphFin, YouTubeReddit-Large). We have
7 reported the corresponding experiments and detailed discussions in the updated paper. The eBay
8 datasets are a collection of the user transactions on **eBay's e-commerce platform**. We thank our
9 industrial collaborator for sharing their datasets in our research. Considering user privacy and security,
10 eBay datasets could only be shared among collaborators. Any researchers who are interested in the
11 eBay datasets, please email our team (jonnyhuanghnu@gmail.com).

12 For easy access, all datasets have been hosted on the open-source platform zenodo ([https://](https://zenodo.org/)
13 zenodo.org/) with a Digital Object Identifier (DOI) 10.5281/zenodo.8267846 ([https://zenodo.](https://zenodo.org/record/8267846)
14 [org/record/8267846](https://zenodo.org/record/8267846)).

15 We have clarified that the difference between BenchTeMP and existing libraries. Furthermore, We
16 illustrate how users can utilize BenchTeMP to test their own temporal GNNs. BenchTeMP can be
17 used for discrete-time dynamic GNNs as well.

18 We provide our response to each individual comment below:

Submitted to the 37th Conference on Neural Information Processing Systems (NeurIPS 2023) Track on Datasets and Benchmarks. Do not distribute.

Comment 1

W1. My major concern is that the adopted datasets are relatively small-scaled, with a maximum number of nodes no more than 100,000. This limitation may impact the potential impact of the benchmark.

19

Response:

21 We appreciate the reviewer for the suggestions! We have included new datasets with up to sev-
22 eral million edges and nodes. We have added *six* datasets (eBay-Small, eBay-Large, Taobao-
23 Large, DGraphFin, YouTubeReddit-Small, YouTubeReddit-Large), including *four large-scale*
24 datasets (eBay-Large, Taobao-Large, DGraphFin, YouTubeReddit-Large). The statistics of the
25 new datasets are shown in Table 1. For easy access, all datasets have been hosted on the open-source
26 platform zenodo (<https://zenodo.org/>) with a Digital Object Identifier (DOI) 10.5281/zen-
27 odo.8267846 (<https://zenodo.org/record/8267846>). In our paper, we present BENCHTEMP,
28 a general benchmark for evaluating temporal graph neural network (TGNN) models over a wide
29 range of tasks and settings. We extensively compare representative TGNN models on the benchmark
30 datasets, regarding different tasks, settings, metrics, and especially model efficiency - **inference time**.
31

- 32 • **eBay-Small** is a subset of the eBay-Large dataset. We sample 38,427 nodes and 384,677 edges
33 from eBay-Large graph according to edge timestamps.
- 34 • **YouTubeReddit-Small** is a collection of massive visual contents on YouTube and long-term
35 community activity on Reddit. This dataset covers a **3**-month period from January to March 2020.
36 Each row in the dataset represents a YouTube video v_i being shared in a subreddit s_j by some
37 user u_k at time t [1]. Nodes are YouTube videos and subreddits, edges are the users' interactions
38 between videos and subreddits. This dynamic graph has 264,443 nodes and 297,732 edges.
- 39 • **eBay-Large** is a million-scale dataset consisting of 1.3 million nodes and 1.1 million edges, which
40 comprises the selected transaction records from the eBay e-commerce platform over a two-month
41 period. eBay-Large is modeled as a user-item graph, where items are heterogeneous entities which
42 include information such as phone numbers, addresses, and email addresses associated with a
43 transaction. We select one month of transactions as seed nodes and then expand each seed node
44 two hops back in time to enrich the topology while maintaining consistency in the distribution of
45 seed nodes.
- 46 • **DGraphFin** is a collection of large-scale dynamic graph datasets, consisting of interactive objects,
47 events and labels that evolve with time. It is a directed, unweighted dynamic graph consisting
48 of millions of nodes and edges, representing a realistic user-to-user social network in financial
49 industry. Nodes are users, and an edge from one user to another means that the user regards the
50 other user as the emergency contact person [2].
- 51 • **Youtube-Reddit-Large** dataset covers **54** months of YouTube video propagation history from
52 January 2018 to June 2022 [1]. This dataset has 5,724,111 nodes and 4,228,523 edges.

Table 1: Dataset statistics of the new datasets.

	<i>Domain</i>	<i># Nodes</i>	<i># Edges</i>
eBay-Small	E-commerce	38,427	384,677
YouTubeReddit-Small [1]	Social	264,443	297,732
eBay-Large	E-commerce	1,333,594	1,119,454
DGraphFin [2]	E-commerce	3,700,550	4,300,999
Youtube-Reddit-Large [1]	Social	5,724,111	4,228,523
Taobao-Large [3, 4]	E-commerce	1,630,453	5,008,745

Table 2: ROC AUC results of new datasets on the *dynamic link prediction task*. The best and second-best results are highlighted as **bold red** and underlined blue. We do not highlight the second-best if the gap is > 0.05 compared with the best result.

Model \ Dataset	Transductive						
	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT
eBay-Small	0.9946 \pm 0.0002	0.9941 \pm 0.0006	0.9984 \pm 0.0003	0.9838 \pm 0.0006	0.9985 \pm 0.0	0.9991 \pm 0.0	<u>0.9978 \pm 0.0003</u>
YouTubeReddit-Small	<u>0.8519 \pm 0.0007</u>	0.8499 \pm 0.0012	0.8432 \pm 0.0032	0.8441 \pm 0.0014	0.7586 \pm 0.0031	0.9003 \pm 0.0031	0.8259 \pm 0.005
eBay-Large	0.9614 \pm 0.0	0.9619 \pm 0.0001	<u>0.9642 \pm 0.0003</u>	0.5311 \pm 0.0003	0.9442 \pm 0.0003	0.9608 \pm 0.0	0.9658 \pm 0.0002
DGraphFin	0.8165 \pm 0.0024	0.8171 \pm 0.0016	0.8683 \pm 0.0023	0.6112 \pm 0.0165	0.5466 \pm 0.0103	<u>0.8611 \pm 0.0035</u>	0.8258 \pm 0.0001
Youtube-Reddit-Large	0.8532 \pm 0.0003	0.8529 \pm 0.0006	0.8458 \pm 0.0025	0.8536 \pm 0.0026	0.7466 \pm 0.0012	0.916 \pm 0.0025	<u>0.8605 \pm 0.0009</u>
Taobao-Large	0.7726 \pm 0.0005	0.7724 \pm 0.001	<u>0.8464 \pm 0.0008</u>	0.5567 \pm 0.0047	0.7771 \pm 0.0068	0.859 \pm 0.0091	0.8188 \pm 0.001
Model \ Dataset	Inductive						
	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT
eBay-Small	0.9696 \pm 0.0007	0.9674 \pm 0.0018	0.9913 \pm 0.0004	0.9698 \pm 0.0006	0.9964 \pm 0.0001	<u>0.9982 \pm 0.0</u>	0.9998 \pm 0.0001
YouTubeReddit-Small	0.7582 \pm 0.0003	0.7545 \pm 0.0009	0.7276 \pm 0.0033	0.7436 \pm 0.0006	0.7533 \pm 0.0016	0.8978 \pm 0.0032	0.9876 \pm 0.0049
eBay-Large	0.7536 \pm 0.0014	0.7515 \pm 0.0006	0.7657 \pm 0.0026	0.5224 \pm 0.0003	0.9459 \pm 0.0001	<u>0.9608 \pm 0.0</u>	0.9999 \pm 0.0001
DGraphFin	0.6884 \pm 0.0051	0.6876 \pm 0.001	0.6439 \pm 0.0089	0.5677 \pm 0.0184	0.5479 \pm 0.009	0.8635 \pm 0.0021	<u>0.7955 \pm 0.0201</u>
Youtube-Reddit-Large	0.7539 \pm 0.0005	0.7554 \pm 0.0003	0.7243 \pm 0.0016	0.7501 \pm 0.0019	0.7327 \pm 0.0016	<u>0.9128 \pm 0.0031</u>	0.9863 \pm 0.006
Taobao-Large	0.7075 \pm 0.0009	0.7042 \pm 0.0006	0.6812 \pm 0.0032	0.5222 \pm 0.0041	0.7787 \pm 0.0103	<u>0.869 \pm 0.010</u>	0.9933 \pm 0.0008
Model \ Dataset	Inductive New-Old						
	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT
eBay-Small	0.9862 \pm 0.0003	0.9836 \pm 0.0016	0.9947 \pm 0.0009	0.9712 \pm 0.002	0.9985 \pm 0.0	<u>0.9988 \pm 0.0</u>	0.9999 \pm 0.0
YouTubeReddit-Small	0.7695 \pm 0.001	0.7655 \pm 0.0018	0.7396 \pm 0.0034	0.7242 \pm 0.0004	0.7573 \pm 0.0022	<u>0.922 \pm 0.0002</u>	0.9967 \pm 0.0014
eBay-Large	0.6109 \pm 0.0244	0.5906 \pm 0.0087	0.8134 \pm 0.0105	0.6363 \pm 0.0605	<u>0.9569 \pm 0.0007</u>	0.8973 \pm 0.0	1.0 \pm 0.0
DGraphFin	0.5768 \pm 0.0071	0.5735 \pm 0.0007	0.5564 \pm 0.0021	0.5742 \pm 0.013	0.5646 \pm 0.0244	<u>0.7702 \pm 0.0043</u>	0.8693 \pm 0.0066
Youtube-Reddit-Large	0.7844 \pm 0.0015	0.7894 \pm 0.0017	0.7623 \pm 0.0031	0.7457 \pm 0.0062	0.7511 \pm 0.0022	<u>0.9356 \pm 0.0004</u>	0.9958 \pm 0.0025
Taobao-Large	0.7023 \pm 0.0015	0.6953 \pm 0.0022	0.6771 \pm 0.0055	0.5104 \pm 0.0106	0.7674 \pm 0.005	<u>0.8458 \pm 0.0043</u>	0.9965 \pm 0.0005
Model \ Dataset	Inductive New-New						
	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT
eBay-Small	0.9388 \pm 0.0009	0.9366 \pm 0.0037	0.9838 \pm 0.0007	0.9556 \pm 0.0007	0.9937 \pm 0.0	<u>0.9975 \pm 0.0</u>	0.9997 \pm 0.0004
YouTubeReddit-Small	0.7436 \pm 0.0015	0.7436 \pm 0.0018	0.7265 \pm 0.0055	0.749 \pm 0.0011	0.7479 \pm 0.004	<u>0.864 \pm 0.0071</u>	0.9868 \pm 0.0049
eBay-Large	0.7526 \pm 0.0013	0.7500 \pm 0.0005	0.7639 \pm 0.0027	0.5196 \pm 0.0002	0.9542 \pm 0.0003	<u>0.9615 \pm 0.0</u>	0.9999 \pm 0.0001
DGraphFin	0.7307 \pm 0.0007	0.7323 \pm 0.0002	0.6843 \pm 0.0131	0.5649 \pm 0.0248	0.5417 \pm 0.0099	0.9051 \pm 0.0028	<u>0.7584 \pm 0.0323</u>
Youtube-Reddit-Large	0.6932 \pm 0.0026	0.7022 \pm 0.0007	0.6703 \pm 0.0024	0.7269 \pm 0.0	0.6942 \pm 0.0028	<u>0.8716 \pm 0.0077</u>	0.9796 \pm 0.0103
Taobao-Large	0.7243 \pm 0.0001	0.7247 \pm 0.0001	0.6885 \pm 0.0024	0.5256 \pm 0.0054	0.7922 \pm 0.0118	<u>0.8906 \pm 0.0088</u>	0.9969 \pm 0.0002

• **Taobao-Large** is a collection of the Taobao user behavior dataset intercepted based on the period 8:00 to 18:00 on 26 November 2017 [4]. Nodes are users and items, and edges are behaviors between users and items, such as favor, click, purchase, and add an item to shopping cart. This public dataset has 1,630,453 nodes and 5,008,74 user-item interaction edges.

A Experiments

We conduct extensive experiments on the tasks of *dynamic link prediction* and *dynamic node classification*. The experimental setup is the same as in the paper.

A.1 Link Prediction Task

We run the link prediction task on 7 TGN models and the new datasets under different settings (Transductive, Inductive, Inductive New-Old, and Inductive New-New). The AUC and AP results for each new datasets are shown in Table 2 and Table 3, respectively. For the four large-scale datasets (eBay-Large, Taobao-Large, DGraphFin, YouTubeReddit-Large), we observe the similar results as in the paper. Specifically, NAT and NeurTW achieve the top-2 performance on almost all datasets under transductive and inductive settings.

A.2 Node Classification Task

The eBay-Small and eBay-Large datasets have node labels, so we conduct dynamic node classification experiments on both the eBay-Small and eBay-Large datasets. The AUC results are shown in Table 4. We can observe the similar results as in the paper. NeurTW achieves the best performance on both eBay-Small and eBay-Large datasets. NAT performs poorly on the node classification task.

Table 3: AP results of new datasets on the *dynamic link prediction task*. The best and second-best results are highlighted as **bold red** and underlined blue. We do not highlight the second-best if the gap is > 0.05 compared with the best result.

		Transductive						
Model \ Dataset	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT	
eBay-Small	0.9938 ± 0.0004	0.9936 ± 0.0006	<u>0.9983 ± 0.0003</u>	0.9819 ± 0.0009	0.9981 ± 0.0	0.9991 ± 0.0	0.9975 ± 0.0002	
YouTubeReddit-Small	<u>0.8612 ± 0.0009</u>	0.8594 ± 0.0012	0.8421 ± 0.0041	0.8515 ± 0.0012	0.7625 ± 0.0042	0.9112 ± 0.0021	0.8325 ± 0.0068	
eBay-Large	0.9318 ± 0.0002	0.9322 ± 0.0002	<u>0.9357 ± 0.0006</u>	0.5239 ± 0.0002	0.9144 ± 0.0004	0.9307 ± 0.0	0.9398 ± 0.0004	
DGraphFin	0.7705 ± 0.0009	0.7705 ± 0.0024	<u>0.8571 ± 0.0009</u>	0.6441 ± 0.0123	0.5431 ± 0.0095	0.8637 ± 0.0014	0.7956 ± 0.0012	
Youtube-Reddit-Large	0.8622 ± 0.0007	<u>0.8632 ± 0.0004</u>	0.8476 ± 0.0022	0.8591 ± 0.0026	0.7475 ± 0.0017	0.9222 ± 0.0013	0.8628 ± 0.0015	
Taobao-Large	0.7164 ± 0.0003	0.7142 ± 0.0008	<u>0.844 ± 0.0011</u>	0.5761 ± 0.0023	0.7616 ± 0.0069	0.8568 ± 0.016	0.7904 ± 0.0008	
		Inductive						
eBay-Small	0.9638 ± 0.0007	0.9619 ± 0.0017	0.9898 ± 0.0005	0.9675 ± 0.0007	0.9953 ± 0.0002	<u>0.9982 ± 0.0</u>	0.9998 ± 0.0001	
YouTubeReddit-Small	0.7866 ± 0.0007	0.7833 ± 0.0009	0.7387 ± 0.0069	0.7551 ± 0.0002	0.7568 ± 0.0031	<u>0.9086 ± 0.0022</u>	0.9872 ± 0.0056	
eBay-Large	0.6989 ± 0.0018	0.6973 ± 0.0007	0.7096 ± 0.0030	0.518 ± 0.0002	0.9174 ± 0.0001	<u>0.9308 ± 0.0</u>	0.9999 ± 0.0001	
DGraphFin	0.6563 ± 0.002	0.6567 ± 0.0009	0.624 ± 0.006	0.5866 ± 0.0123	0.5428 ± 0.0082	0.8626 ± 0.0012	<u>0.7053 ± 0.0185</u>	
Youtube-Reddit-Large	0.7796 ± 0.0009	0.7818 ± 0.0009	0.73 ± 0.0029	0.7587 ± 0.0025	0.7353 ± 0.0022	<u>0.9192 ± 0.0022</u>	0.9849 ± 0.0071	
Taobao-Large	0.6763 ± 0.0011	0.6746 ± 0.0011	0.6664 ± 0.0012	0.5315 ± 0.0027	0.7533 ± 0.011	<u>0.8596 ± 0.0205</u>	0.9941 ± 0.0007	
		Inductive New-Old						
eBay-Small	0.9849 ± 0.0007	0.9836 ± 0.0013	0.9931 ± 0.0008	0.9682 ± 0.0028	0.9985 ± 0.0001	<u>0.999 ± 0.0</u>	0.9999 ± 0.0	
YouTubeReddit-Small	0.7963 ± 0.0013	0.7937 ± 0.0014	0.729 ± 0.0086	0.7296 ± 0.0013	0.762 ± 0.0041	<u>0.9244 ± 0.0015</u>	0.9966 ± 0.0016	
eBay-Large	0.5670 ± 0.0186	0.5870 ± 0.0074	0.8024 ± 0.0060	0.6504 ± 0.0385	<u>0.9592 ± 0.0008</u>	0.8458 ± 0.0	1.0 ± 0.0	
DGraphFin	0.6005 ± 0.0048	0.5872 ± 0.0059	0.5753 ± 0.0062	0.5927 ± 0.0058	0.5669 ± 0.0269	<u>0.7572 ± 0.0025</u>	0.8184 ± 0.0088	
Youtube-Reddit-Large	0.808 ± 0.0014	0.8142 ± 0.0019	0.7472 ± 0.0043	0.7526 ± 0.0097	0.7553 ± 0.0025	<u>0.9368 ± 0.0009</u>	0.9953 ± 0.0028	
Taobao-Large	0.7009 ± 0.0013	0.698 ± 0.0014	0.6879 ± 0.0008	0.5254 ± 0.0074	0.7597 ± 0.0053	<u>0.8459 ± 0.0103</u>	0.9969 ± 0.0004	
		Inductive New-New						
eBay-Small	0.923 ± 0.001	0.9226 ± 0.0024	0.98 ± 0.0007	0.9505 ± 0.0009	0.991 ± 0.0001	<u>0.9973 ± 0.0</u>	0.9997 ± 0.0004	
YouTubeReddit-Small	0.7578 ± 0.0015	0.7582 ± 0.0021	0.7564 ± 0.0043	0.7718 ± 0.0023	0.7498 ± 0.004	<u>0.8868 ± 0.0034</u>	0.9861 ± 0.0063	
eBay-Large	0.6976 ± 0.0016	0.6957 ± 0.0007	0.7078 ± 0.0031	0.5154 ± 0.0001	0.93 ± 0.0003	<u>0.9318 ± 0.0</u>	0.9999 ± 0.0001	
DGraphFin	0.6802 ± 0.0005	0.6811 ± 0.0002	0.6526 ± 0.0098	0.5831 ± 0.0184	0.5379 ± 0.0071	0.8977 ± 0.0014	0.6529 ± 0.0249	
Youtube-Reddit-Large	0.7038 ± 0.0024	0.7115 ± 0.0007	0.6979 ± 0.002	0.7414 ± 0.0012	0.6965 ± 0.004	<u>0.8848 ± 0.0023</u>	0.9761 ± 0.0134	
Taobao-Large	0.6738 ± 0.0005	0.6742 ± 0.0005	0.6611 ± 0.0011	0.53 ± 0.0023	0.7521 ± 0.0127	<u>0.8738 ± 0.0145</u>	0.9973 ± 0.0001	

Table 4: ROC AUC results for the *dynamic node classification task* on the eBay datasets. The top-2 results are highlighted as **bold red** and underlined blue.

Model \ Dataset	JODIE	DyRep	TGN	TGAT	CAWN	NeurTW	NAT
eBay-Small	0.9274 \pm 0.0017	0.8677 \pm 0.0356	0.913 \pm 0.0025	<u>0.9342 \pm 0.0002</u>	0.9305 \pm 0.0001	0.9529 \pm 0.0002	0.6797 \pm 0.0115
eBay-Large	0.7244 \pm 0.0002	0.7246 \pm 0.0	0.6586 \pm 0.0129	0.672 \pm 0.0016	<u>0.7710 \pm 0.0002</u>	0.7859 \pm 0.0	0.5304 \pm 0.0011

72 A.3 Efficiency

73 Considering many real world applications and , we add **the inference time** metric to evaluate the
74 efficiency of models. The inference time comparison per 100,000 edges is shown in Figure 1.
75 According to the figure, we can observe the similar model efficiency results as in the paper. In
76 terms of the inference time, JODIE, DyRep, TGN and TGAT are faster, while CAWN and NeurTW
77 are much slower. NAT is relatively faster than temporal walk-based methods through caching and
78 parallelism optimizations, *achieving a good trade-off between model quality and efficiency.*

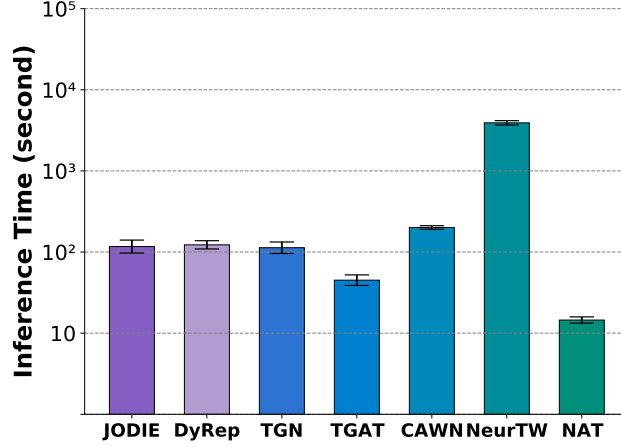


Figure 1: Inference time comparison per 100,000 edges.

Comment 2

W2. The paper claims to contribute to the unification of pipelines for existing temporal GNNs, but it is worth noting that there are already libraries available, such as PyG Temporal and TGL (VLDB’22), which address this aspect, albeit without providing benchmarks. It would be helpful to clarify how the proposed method differs from these libraries or whether they can be used together.

79

Response:

81 We appreciate this suggestion! We have carefully through your comments and compare BenchTeMP
82 with existing libraries or methods, such as PyG Temporal [5] and TGL [6].

83 PyG Temporal [5] is a temporal graph neural network extension library for PyTorch Geometric.
84 PyG Temporal consists of state-of-the-art deep learning and parametric learning methods to process
85 spatio-temporal signals. PyG Temporal provides constant time difference graph neural networks on
86 dynamic and static graphs. Users can build their own models based on PyG Temporal. BenchTeMP
87 proposed in this paper is a benchmark library, which consists of *benchmark dataset*, *DataLoader*,
88 *EdgeSampler*, *Evaluator*, *EarlyStopMonitor*, and *Leaderboard*. For evaluating TGNN models onto
89 the same ground and compares them comprehensively, BenchTeMP provided a unified benchmark
90 pipeline shown in Figure 4 in the paper for users to evaluate their own models constructed by PyG
91 Temporal [5]. Of course, BenchTeMP and PyG Temporal can be used together. For example, users
92 can construct their own models by PyG Temporal framework, and evaluate the performance of models
93 by BenchTeMP.

94 TGL [6] comprises five main components, *a temporal sampler*, *a mailbox*, *a node memory module*,
95 *a memory updater*, and *a message passing engine*. Considering the memory module and memory
96 updater components, TGL [6] is a framework for memory-based TGNN models, such as JODIE [7],
97 TGAT [8], TGN [9], and APAN [10]. However, BenchTeMP is a general benchmark framework for
98 evaluating, no matter memory-based TGNNs (JODIE, DyRep, TGAT, TGN), TGNNs based on motifs
99 (CAWN[11], NeurTW [3]), or even TGNNs based on joint-neighborhood operation (NAT [12]).
100 Furthermore, BenchTeMP provides diverse workloads for evaluating TGNNs, regarding different
101 tasks (dynamic link prediction and dynamic node classification), settings (transductive, inductive,
102 New-Old, and New-New), metrics, and efficiency (runtime, running memory, inference time).

Comment 3

W3. The benchmark code page could benefit from improved documentation. For example, it is not immediately clear how users can utilize the benchmark to test their own temporal GNNs.

Response:

Thanks for your valuable suggestion! We have updated the benchmark code page (<https://github.com/qianghuangwhu/benchtemp>) and add a Section - *Usage Example*. In future work, we will continuously update our benchmark code page for users. Contributions and issues from the community are eagerly welcomed, with which we can together push forward the TGNN research.

Comment 4

W4. While the main focus of the paper is on temporal GNNs, it would be valuable to discuss whether the proposed benchmark can be used for discrete-time dynamic GNNs as well.

Response:

Thanks for this valuable comment! Of course, BenchTeMP proposed by us can be used for discrete-time dynamic GNNs as well!

For example, DySAT [13] is a discrete-time dynamic graph model [9]. DySAT treats temporal interactions as graph snapshots. The input of DySAT are snapshot at each time step, and the output of DySAT model is node representations at each time step. For link prediction task, Using the node embeddings trained on graph snapshots up to time step t , single-step link prediction predicts the connections between nodes at time step $t + 1$. Exactly the same as the RandEdgeSampler of BenchTeMP, DySAT also performs negative sampling operation.

For example, the code of negative sampling in DySAT as follows, See file <https://github.com/aravindsankar28/DySAT/blob/master/utils/preprocess.py> for code details:

```
121 1 # python3
122 2 # Create train edges.
123 3 train_edges_false = []
124 4 while len(train_edges_false) < len(train_edges):
125 5     idx_i = np.random.randint(0, adj.shape[0])
126 6     idx_j = np.random.randint(0, adj.shape[0])
127 7     if idx_i == idx_j:
128 8         continue
129 9     if ismember([idx_i, idx_j], edges_all):
13010         continue
13111     if ismember([idx_j, idx_i], edges_all):
13212         continue
13313     if train_edges_false:
13414         if ismember([idx_j, idx_i], np.array(train_edges_false)):
13515             continue
13616         if ismember([idx_i, idx_j], np.array(train_edges_false)):
13717             continue
13818     train_edges_false.append([idx_i, idx_j])
```

The code of RandEdgeSampler in BenchTeMP below, See file (<https://github.com/qianghuangwhu/benchtemp/blob/master/lp/edgesampler.py>) for code details:

```
141 1 # BenchTeMP RandEdgeSampler
142 2 class RandEdgeSampler:
143 3     def __init__(self, src_list, dst_list, seed=None):
```

```

144 4         self.seed = None
145 5         self.src_list = np.unique(src_list)
146 6         self.dst_list = np.unique(dst_list)
147 7
148 8         if seed is not None:
149 9             self.seed = seed
150 10            self.random_state = np.random.RandomState(self.seed)
151 11
152 12     def sample(self, size):
153 13         if self.seed is None:
154 14             src_index = np.random.randint(0, len(self.src_list), size)
155 15             dst_index = np.random.randint(0, len(self.dst_list), size)
156 16         else:
157 17
158 18             src_index = self.random_state.randint(0, len(self.src_list
159 19             ), size)
160 20             dst_index = self.random_state.randint(0, len(self.dst_list
161 21             ), size)
162 22             return self.src_list[src_index], self.dst_list[dst_index]
163 23
164 24     def reset_random_state(self):
165 25         self.random_state = np.random.RandomState(self.seed)

```

166 Besides, the data format of datasets and data loading operation of DySAT are exactly the same as
167 BenchTeMP.

168 Therefore, by the above analysis, we can conclude that BenchTeMP proposed by us can be used
169 for discrete-time dynamic graph models as well.

References

- [1] Yiqiao Jin, Yeon-Chang Lee, Kartik Sharma, Meng Ye, Karan Sikka, Ajay Divakaran, and Srijan Kumar. Predicting information pathways across online communities. *arXiv preprint arXiv:2306.02259*, 2023.
- [2] Xuanwen Huang, Yang Yang, Yang Wang, Chunping Wang, Zhisheng Zhang, Jiarong Xu, Lei Chen, and Michalis Vazirgiannis. Dgraph: A large-scale financial dataset for graph anomaly detection. *Advances in Neural Information Processing Systems*, 35:22765–22777, 2022.
- [3] Ming Jin, Yuan-Fang Li, and Shirui Pan. Neural temporal walks: Motif-aware representation learning on continuous-time dynamic graphs. In *Advances in Neural Information Processing Systems*, 2022.
- [4] Han Zhu, Xiang Li, Pengye Zhang, Guozheng Li, Jie He, Han Li, and Kun Gai. Learning tree-based deep model for recommender systems. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1079–1088, 2018.
- [5] Benedek Rozemberczki, Paul Scherer, Yixuan He, George Panagopoulos, Alexander Riedel, Maria Astefanoaei, Oliver Kiss, Ferenc Beres, , Guzman Lopez, Nicolas Collignon, and Rik Sarkar. PyTorch Geometric Temporal: Spatiotemporal Signal Processing with Neural Machine Learning Models. In *Proceedings of the 30th ACM International Conference on Information and Knowledge Management*, page 4564–4573, 2021.
- [6] Hongkuan Zhou, Da Zheng, Israt Nisa, Vasileios Ioannidis, Xiang Song, and George Karypis. Tgl: a general framework for temporal gnn training on billion-scale graphs. *Proceedings of the VLDB Endowment*, 15(8):1572–1580, 2022.
- [7] Srijan Kumar, Xikun Zhang, and Jure Leskovec. Predicting dynamic embedding trajectory in temporal interaction networks. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1269–1278, 2019.
- [8] Da Xu, Chuanwei Ruan, Evren Korpeoglu, Sushant Kumar, and Kannan Achan. Inductive representation learning on temporal graphs. In *International Conference on Learning Representations*, 2020.
- [9] Emanuele Rossi, Ben Chamberlain, Fabrizio Frasca, Davide Eynard, Federico Monti, and Michael Bronstein. Temporal graph networks for deep learning on dynamic graphs. *arXiv preprint arXiv:2006.10637*, 2020.
- [10] Xuhong Wang, Ding Lyu, Mengjian Li, Yang Xia, Qi Yang, Xinwen Wang, Xinguang Wang, Ping Cui, Yupu Yang, Bowen Sun, et al. Apan: Asynchronous propagation attention network for real-time temporal graph embedding. In *Proceedings of the 2021 international conference on management of data*, pages 2628–2638, 2021.
- [11] Yanbang Wang, Yen-Yu Chang, Yunyu Liu, Jure Leskovec, and Pan Li. Inductive representation learning in temporal networks via causal anonymous walks. In *International Conference on Learning Representations*, 2021.
- [12] Yuhong Luo and Pan Li. Neighborhood-aware scalable temporal network representation learning. In *The First Learning on Graphs Conference*, 2022.
- [13] Aravind Sankar, Yanhong Wu, Liang Gou, Wei Zhang, and Hao Yang. Dysat: Deep neural representation learning on dynamic graphs via self-attention networks. In *Proceedings of the 13th international conference on web search and data mining*, pages 519–527, 2020.